

Accessing Intel ICH/PCH GPIOs

2017-11-08

Tags: [software](#), [electronics](#).

In this note I describe the method required for accessing GPIOs on (in theory) every Intel chipset from ICH0 to (at least) Series 100 PCH.

Overview

Historically (going as far back as ICH0), Intel used a special PCI function to implement a multitude of legacy functions (such as `A20#`, the [LPC bus](#), etc). The Intel lingo for this is “D31:F0”, as in “device 31, function 0”, but it shows up in `lspci` as `00:1f.0` ISA bridge. In this note, I will use the Intel terminology for PCI devices.

Table of Contents

[Overview](#)

- [GPIOs on ICH0..9 and Series 5..9 PCH](#)
- [GPIOs on Series 100 PCH](#)

[Intel documentation](#)

[Code](#)

GPIOs on ICH0..9 and Series 5..9 PCH

On these platforms, D31:F0 has an I/O BAR. The size of the I/O space was increased for ICH6, but otherwise the GPIO interface is identical. It is comprised of a set of I/O ports. To access them, it is enough to place the I/O space of the function in the global I/O space, and enable decoding.

The implementation is as follows:

- The vendor and device IDs are checked to make sure we're not crashing the system.
- The I/O BAR is configured via `libpci`.
- The I/O permissions are requested via `ioperm`.
- Finally, GPIO registers can be used.

GPIOs on Series 100 PCH

On these platforms, D31:F0 is solely dedicated to being an LPC bridge. The GPIOs are located in what Intel calls “private configuration space”, accessible through a “primary to sideband bridge” through “target port IDs”. All this seems extremely opaque, but in reality very little has changed.

The “primary to sideband bridge” is simply a PCI function (located at D31:F1) that has BAR0, a memory BAR, “private configuration space”, initialized by platform firmware to point to some location it finds convenient. To prevent the OS from reassigning the BAR, the firmware “hides” the device, namely sets a bit in the configuration space that causes all reads to return all-ones. (Writes still go through.) The “target port ID” is the bits [23:16] of the D31:F1 BAR0.

The implementation is as follows:

- The vendor and device IDs are checked to make sure we're not crashing the system.
- The D31:F1 function is blindly enabled bypassing the operating system as to avoid changing the BAR, and verified to have been enabled correctly.
- The memory BAR is read from D31:F1.
- The D31:F1 function is quickly disabled again.
- The address read from BAR is mapped into the process' address space.
- Finally, GPIO registers can be used.

Intel documentation

- Intel® 82801EB I/O Controller Hub 5 (ICH5) / Intel® 82801ER I/O Controller Hub 5 R (ICH5R) Datasheet (Document [252516-001](#))
- Intel® 9 Series Chipset Family Platform Controller Hub (PCH) Datasheet (Document [330550-002](#))
- Intel® 100 Series and Intel® C230 Series Chipset Family Platform Controller Hub (PCH) Datasheet - Volume 1 of 2 (Document [332690-004EN](#))
- Intel® 100 Series and Intel® C230 Series Chipset Family Platform Controller Hub (PCH) Datasheet - Volume 2 of 2 (Document [332691-002EN](#))

Code

Run `make && sudo ./gpioke` and enjoy a printout of GPIO status live from your chipset. This demo was written very carefully and is not supposed to ever crash your machine. However, it has not undergone a lot of live testing.

Makefile

```

1  CFLAGS = -std=c11 -Wall -g
2  LIBS   = -lpci
3
4  gpioke: gpioke.c
5          $(CC) $(CFLAGS) -o $@ $^ $(LIBS)

```

gpioke.c ([download](#))

```

1  /*
2  * See:
3  * - http://lxr.free-electrons.com/source/drivers/mfd/lpc_ich.c
4  * - http://lxr.free-electrons.com/source/drivers/gpio/gpio-ich.c
5  * - Intel document 252516-001 (ICH5)
6  * - Intel document 330550-002 (9 Series PCH)
7  * - Intel documents 332690-004 and 332691-002EN (100 Series PCH)
8  */
9
10 #include <stdio.h>
11 #include <stdint.h>
12 #include <unistd.h>
13 #include <fcntl.h>

```

```

14 #include <pci/pci.h>
15 #include <sys/io.h>
16 #include <sys/mman.h>
17 #include <sys/errno.h>
18
19 /*
20  * GPIO register offsets in GPIO I/O space.
21  * Each chunk of 32 GPIOs is manipulated via its own USE_SELx, IO_SELx, and
22  * LVLx registers. Logic in the read/write functions takes a register and
23  * an absolute bit number and determines the proper register offset and bit
24  * number in that register. For example, to read the value of GPIO bit 50
25  * the code would access offset ichx_regs[2(=GPIO_LVL)][1(=50/32)],
26  * bit 18 (50%32).
27  */
28 enum GPIO_REG {
29     GPIO_USE_SEL = 0,
30     GPIO_IO_SEL,
31     GPIO_LVL,
32 };
33
34 static const uint8_t ichx_regs[4][3] = {
35     {0x00, 0x30, 0x40}, /* USE_SEL[1-3] offsets */
36     {0x04, 0x34, 0x44}, /* IO_SEL[1-3] offsets */
37     {0x0c, 0x38, 0x48}, /* LVL[1-3] offsets */
38 };
39
40 /*
41  * Generic PCI configuration space registers.
42  */
43 #define REG_VENDOR      0x00
44 #define REG_DEVICE      0x04
45
46 /*
47  * D31:F0 configuration space registers.
48  */
49 #define REG_ICH0_GPIOBASE 0x58
50 #define REG_ICH0_GC       0x5c
51
52 #define REG_ICH6_GPIOBASE 0x48
53 #define REG_ICH6_GC       0x4c
54
55 #define REG_ICHx_GC_EN    0x10
56 #define REG_ICHx_GC_GLE   0x01
57
58 /*
59  * D31:F1 configuration space registers.
60  */
61 #define REG_P2SB_BAR      0x10
62 #define REG_P2SB_BARH     0x14
63 #define REG_P2SB_CTRL     0xe0
64
65 #define REG_P2SB_CTRL_HIDE 0x0100
66
67 /*
68  * P2SB private registers.
69  */
70 #define P2SB_PORTID_SHIFT 16
71 #define P2SB_PORT_GPIO3   0xAC
72 #define P2SB_PORT_GPIO2   0xAD
73 #define P2SB_PORT_GPIO1   0xAE
74 #define P2SB_PORT_GPIO0   0xAF
75
76 /*
77  * GPIO sideband registers.
78  */
79 #define REG_PCH_GPIO_FAMBAR 0x8
80 #define REG_PCH_GPIO_PADBAR 0xc
81
82 #define REG_PCH_GPIO_PAD_OWN 0x20

```

```

83 #define REG_PCH_GPIO_HOSTSW_OWN    0xd0
84 #define REG_PCH_GPIO_GPI_IS       0x100
85 #define REG_PCH_GPIO_GPI_IE       0x120
86 #define REG_PCH_GPIO_GPE_STS      0x140
87 #define REG_PCH_GPIO_GPE_EN       0x160
88 #define REG_PCH_GPIO_SMI_STS      0x184
89 #define REG_PCH_GPIO_SMI_EN       0x1a4
90 #define REG_PCH_GPIO_NMI_STS      0x1c4
91 #define REG_PCH_GPIO_NMI_EN       0x1e4
92
93 #define REG_PCH_GPIO_DW0_PMODE     0x1600
94 #define REG_PCH_GPIO_DW0_RXDIS     0x0200
95 #define REG_PCH_GPIO_DW0_TXDIS     0x0100
96 #define REG_PCH_GPIO_DW0_RXSTATE   0x0002
97 #define REG_PCH_GPIO_DW0_TXSTATE   0x0001
98
99 #define REG_PCH_GPIO_DW1_TERM_NONE 0x0
100 #define REG_PCH_GPIO_DW1_TERM_5K_PD 0x2
101 #define REG_PCH_GPIO_DW1_TERM_20K_PD 0x4
102 #define REG_PCH_GPIO_DW1_TERM_5K_PU 0xa
103 #define REG_PCH_GPIO_DW1_TERM_20K_PU 0xc
104 #define REG_PCH_GPIO_DW1_TERM_NATIVE 0xf
105
106 /*
107  * Helper functions.
108  */
109
110 #define MSG(...) do { \
111     fprintf(stderr, "[%s] " __VA_ARGS__); fprintf(stderr, "\n"); \
112 } while(0)
113 #define ERR(...) do { \
114     fprintf(stderr, "[-] " __VA_ARGS__); fprintf(stderr, "\n"); \
115     return 1; \
116 } while(0)
117 #define DIE(...) do { *fatal = 1; ERR(__VA_ARGS__) } while(0)
118
119 struct pci_dev *pci_find_dev(struct pci_access *pci, uint8_t bus, uint8_t dev, ui
120     for(struct pci_dev *it = pci->devices; it; it = it->next) {
121         if(it->bus == bus && it->dev == dev && it->func == func) return it;
122     }
123     return NULL;
124 }
125
126 /*
127  * Finally, our main logic!
128  */
129
130 int try_ich(struct pci_access *pci,
131             uint16_t reg_gpiobase, uint16_t reg_gc,
132             const char *desc, int *fatal) {
133     MSG("Checking for a %s system", desc);
134
135     struct pci_dev *d31f0 = pci_find_dev(pci, 0, 31, 0);
136     uint32_t gpiobase = pci_read_long(d31f0, reg_gpiobase);
137     uint8_t gc = pci_read_byte(d31f0, reg_gc);
138     MSG("GPIOBASE=%08x, GC=%02x", gpiobase, gc);
139
140     if(gpiobase == 0xffffffff) {
141         *fatal = 1;
142         ERR("Cannot read GPIOBASE, are you running me as root?");
143     } else if(gpiobase == 0) {
144         ERR("GPIOBASE not implemented at %04x", reg_gpiobase);
145     } else if(!(gpiobase & 1)) {
146         *fatal = 1;
147         ERR("GPIOBASE is not an I/O BAR");
148     }
149
150     if(!(gpiobase & 0xffffc)) {
151         const uint32_t DEFAULT_GPIOBASE = 0x0480;

```

```

152
153     MSG("GPIOBASE is not configured, setting to %08x and hoping this works", DEFA
154     pci_write_long(d31f0, reg_gpiobase, DEFAULT_GPIOBASE);
155     gpiobase = pci_read_long(d31f0, reg_gpiobase);
156     if((gpiobase & 0xffffc) != DEFAULT_GPIOBASE) {
157         ERR("Cannot set GPIOBASE");
158     }
159 }
160
161 MSG("GPIO decoding is %s", (gc & REG_ICHx_GC_EN) ? "enabled" : "disabled");
162 MSG("GPIO lockdown is %s", (gc & REG_ICHx_GC_GLE) ? "enabled" : "disabled");
163
164 if(!(gc & REG_ICHx_GC_EN)) {
165     MSG("Enabling GPIO decoding");
166     pci_write_byte(d31f0, reg_gc, gc | REG_ICHx_GC_EN);
167     gc = pci_read_byte(d31f0, reg_gc);
168     if(!(gc & REG_ICHx_GC_EN)) {
169         ERR("Cannot enable GPIO decoding");
170     }
171 }
172
173 gpiobase &= 0xffffc;
174 if(ioperm(gpiobase, 128, 1) == -1) {
175     ERR("Cannot access I/O ports %04x:%04x", gpiobase, gpiobase + 128);
176 }
177
178 for(int n = 1; n < 3; n++) {
179     MSG("USE_SEL%d=%08x", n, inl(gpiobase + ichx_regs[GPIO_USE_SEL][n]));
180     MSG("IO_SEL%d=%08x", n, inl(gpiobase + ichx_regs[GPIO_IO_SEL][n]));
181     MSG("LVL%d=%08x", n, inl(gpiobase + ichx_regs[GPIO_LVL][n]));
182 }
183
184 return 0;
185 }
186
187 int get_pch_sbreg_addr(struct pci_access *pci, pciaddr_t *sbreg_addr) {
188     MSG("Checking for a Series 10 PCH system");
189
190     struct pci_dev *d31f1 = pci_get_dev(pci, 0, 0, 31, 1);
191     pci_fill_info(d31f1, PCI_FILL_IDENT);
192     if(d31f1->vendor_id == 0xffff) {
193         MSG("Cannot find D31:F1, assuming it is hidden by firmware");
194
195         uint32_t p2sb_ctrl = pci_read_long(d31f1, REG_P2SB_CTRL);
196         MSG("P2SB_CTRL=%02x", p2sb_ctrl);
197         if(!(p2sb_ctrl & REG_P2SB_CTRL_HIDE)) {
198             ERR("D31:F1 is hidden but P2SB_E1 is not 0xff, bailing out");
199         }
200
201         MSG("Unhiding P2SB");
202         pci_write_long(d31f1, REG_P2SB_CTRL, p2sb_ctrl & ~REG_P2SB_CTRL_HIDE);
203
204         p2sb_ctrl = pci_read_long(d31f1, REG_P2SB_CTRL);
205         MSG("P2SB_CTRL=%02x", p2sb_ctrl);
206         if(p2sb_ctrl & REG_P2SB_CTRL_HIDE) {
207             ERR("Cannot unhide PS2B");
208         }
209
210         pci_fill_info(d31f1, PCI_FILL_RESCAN | PCI_FILL_IDENT);
211         if(d31f1->vendor_id == 0xffff) {
212             ERR("P2SB unhidden but does not enumerate, bailing out");
213         }
214     }
215
216     pci_fill_info(d31f1, PCI_FILL_RESCAN | PCI_FILL_IDENT | PCI_FILL_BASES);
217     if(d31f1->vendor_id != 0x8086) {
218         ERR("Vendor of D31:F1 is not Intel");
219     } else if((uint32_t)d31f1->base_addr[0] == 0xffffffff) {
220         ERR("SBREG_BAR is not implemented in D31:F1");

```

```

221     }
222
223     *sbreg_addr = d31f1->base_addr[0] &~ 0xf;
224     MSG("SBREG_ADDR=%08lx", *sbreg_addr);
225
226     MSG("Hiding P2SB again");
227     uint32_t p2sb_ctrl = pci_read_long(d31f1, REG_P2SB_CTRL);
228     pci_write_long(d31f1, REG_P2SB_CTRL, p2sb_ctrl | REG_P2SB_CTRL_HIDE);
229
230     pci_fill_info(d31f1, PCI_FILL_RESCAN | PCI_FILL_IDENT);
231     if(d31f1->vendor_id != 0xffff) {
232         ERR("Cannot hide P2SB");
233     }
234
235     return 0;
236 }
237
238 uint32_t sideband_read(void *sbmap, uint8_t port, uint16_t reg) {
239     uintptr_t addr = ((uintptr_t)sbmap + (port << P2SB_PORTID_SHIFT) + reg);
240     return *((volatile uint32_t *)addr);
241 }
242
243 int try_pch(struct pci_access *pci) {
244     pciaddr_t sbreg_addr;
245     if(get_pch_sbreg_addr(pci, &sbreg_addr)) {
246         MSG("Re-enumerating PCI devices will probably crash the system");
247         ERR("Probing Series 100 PCH failed");
248     }
249
250     int memfd = open("/dev/mem", O_RDWR);
251     if(memfd == -1) {
252         ERR("Cannot open /dev/mem");
253     }
254
255     void *sbmap = mmap((void*)sbreg_addr, 1<<24, PROT_READ|PROT_WRITE, MAP_SHARED,
256                       memfd, sbreg_addr);
257     if(sbmap == MAP_FAILED) {
258         if(errno == EPERM) {
259             MSG("Is your kernel configured with CONFIG_DEVMEM_STRICT=n?");
260             MSG("Try rebooting and specifying iomem=relaxed on kernel command line.");
261         }
262         ERR("Cannot map SBREG");
263     }
264
265     close(memfd);
266
267     for(unsigned port = 0; port < 4; port++) {
268         uint16_t port_id = P2SB_PORT_GPIO0 - port;
269         uint32_t padbar = sideband_read(sbmap, port_id, REG_PCH_GPIO_PADBAR);
270         MSG("GPIO%d_PADBAR=%x", port, padbar);
271
272         for(unsigned pad = 0; pad < 32; pad++) {
273             uint32_t dw0 = sideband_read(sbmap, port_id, padbar + pad * 8);
274             uint32_t dw1 = sideband_read(sbmap, port_id, padbar + pad * 8 + 4);
275             if(dw1 == 0) {
276                 // Not documented as such, but appears to be a reliable last pad marker.
277                 break;
278             }
279
280             const char *state = "???", *rxstate = "", *txstate = "";
281             if((dw0 & REG_PCH_GPIO_DW0_PMODE) != 0) {
282                 state = "Native";
283             } else if((dw0 & REG_PCH_GPIO_DW0_TXDIS) != 0 &&
284                      (dw0 & REG_PCH_GPIO_DW0_RXDIS) != 0) {
285                 state = "Off";
286             } else {
287                 state = "GPIO";
288                 if((dw0 & REG_PCH_GPIO_DW0_RXDIS) == 0) {
289                     if((dw0 & REG_PCH_GPIO_DW0_RXSTATE) != 0) {

```

```

290         rxstate = " InHigh";
291     } else {
292         rxstate = " InLow";
293     }
294 }
295
296 if((dw0 & REG_PCH_GPIO_DW0_TXDIS) == 0) {
297     if((dw0 & REG_PCH_GPIO_DW0_TXSTATE) != 0) {
298         txstate = " OutHigh";
299     } else {
300         txstate = " OutLow";
301     }
302 }
303 }
304
305 const char *pull = "???";
306 switch(dw1 >> 10) {
307     case REG_PCH_GPIO_DW1_TERM_NONE:    pull = "None";    break;
308     case REG_PCH_GPIO_DW1_TERM_5K_PD:   pull = "Dn5k";    break;
309     case REG_PCH_GPIO_DW1_TERM_20K_PD:  pull = "Dn20k";   break;
310     case REG_PCH_GPIO_DW1_TERM_5K_PU:   pull = "Up5k";    break;
311     case REG_PCH_GPIO_DW1_TERM_20K_PU:  pull = "Up20k";   break;
312     case REG_PCH_GPIO_DW1_TERM_NATIVE:  pull = "Native";  break;
313 }
314
315 printf("[+] GPIO%d_PAD%d: DW0=%08x DW1=%08x State=%s%s%s Pull=%s\n",
316        port, pad, dw0, dw1, state, rxstate, txstate, pull);
317 }
318 }
319
320 return 0;
321 }
322
323 int create_pci(int method, struct pci_access **pci_out) {
324     struct pci_access *pci = pci_alloc();
325     pci->method = method;
326     pci_init(pci);
327     pci_scan_bus(pci);
328
329     struct pci_dev *d31f0 = pci_find_dev(pci, 0, 31, 0);
330     if(!d31f0) {
331         ERR("Cannot find D31:F0");
332     }
333
334     pci_fill_info(d31f0, PCI_FILL_IDENT | PCI_FILL_BASES);
335     if(d31f0->vendor_id != 0x8086) {
336         ERR("Vendor of D31:F0 is not Intel");
337     }
338
339     *pci_out = pci;
340     return 0;
341 }
342
343 int main() {
344     struct pci_access *pci;
345     if(create_pci(PCI_ACCESS_AUTO, &pci)) {
346         MSG("Is this an Intel platform?");
347         return 1;
348     }
349
350     int fatal = 0;
351     if(try_ich(pci, REG_ICH0_GPIOBASE, REG_ICH0_GC,
352               "ICH0..ICH5", &fatal) && fatal) {
353         return 1;
354     } else if(try_ich(pci, REG_ICH6_GPIOBASE, REG_ICH6_GC,
355                       "ICH6..ICH9 or Series 5..9 PCH", &fatal) && fatal) {
356         return 1;
357     } else {
358         pci_cleanup(pci);

```



```
359
360 // Letting Linux discover P2SB (and reassign its BAR) hangs the system,
361 // so we need to enumerate the device bypassing it.
362 if(create_pci(PCI_ACCESS_I386_TYPE1, &pci)) {
363     return 1;
364 }
365
366 if(try_pch(pci)) {
367     return 1;
368 }
369 }
370
371 printf("[+] Done\n");
372 return 0;
373 }
```

Want to discuss this note? Drop me a [letter](#).

[Drafts](#)

whitequark © 2014-2018 [CC0](#)